

Dynamic Software Updates: A VM-centric Approach

Suriya Subramanian

Department of Computer Sciences
The University of Texas at Austin

July 14, 2009

Motivation

- Software applications change all the time
- Deployed systems must be updated with bug fixes, new features
- Updating typically involves: stop, apply patch, restart








Your request could not be processed at this time




The service is currently unavailable.

Please try again in a few minutes. If this problem persists, report an issue to eBay support.




Your request could not be processed

 The service is currently unavailable.
Please try again in a few minutes. If

The fundamental problem is losing state
because of downtime.


Dynamic software updating




Overheard on programming.reddit



Overheard on programming.reddit



Overheard on programming.reddit



Dynamic updating systems

- Special-purpose architectures, application-specific solutions exist
- General-purpose solutions gaining strength
 - K42, Ksplice for OS updates
 - Polus, Ginseng for C applications
- Not for managed languages

DSU opportunity for managed languages

DSU Solutions for C/C++ typically

- Require special compilation
- Statically/dynamically insert indirection for function calls
- Restrict structure updates, require extra allocation
- Impose space/time overheads on normal execution
- Make type-safety for updates difficult
- Not multi-threaded

Our solution

- Jvolve - a Java Virtual Machine with DSU support
- Key insight: Extend existing VM services
- No DSU-related overhead during normal execution
- Support updates to real world applications

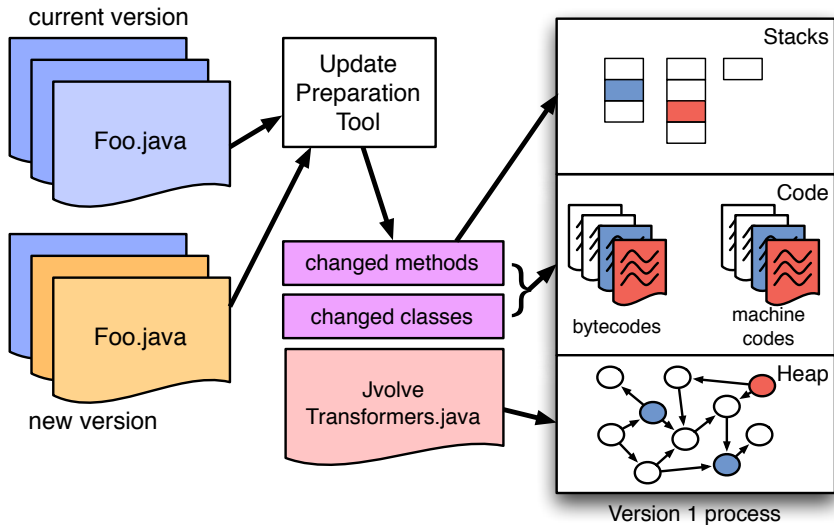
*Dynamic software updating in managed languages can be achieved in a **safe, flexible and efficient** manner by naturally extending existing VM services.*

DSU support should be a standard feature of future VMs.

Outline

- Introduction
- Jvolve
 - VM Implementation
 - Evaluation
- Future Work
- Conclusion

JVOLVE - System overview



Supported updates

- Changes within the body of a method

```
public static void main(String args[]) {  
    System.out.println("Hello, World.");  
+   System.out.println("Hello again, World.");  
}
```

- Class signature updates

- Add, remove, change the type signature of fields and methods

```
public class Line {  
-   private final Point2D p1, p2;  
+   private final Point3D p1, p2;  
    ...  
}
```

- Signature updates require an object transformer function

Example of an update (JavaEmailServer)

```
public class User {
    private final String username, domain, password;
-   private String[] forwardAddresses;
+   private EmailAddress[] forwardAddresses;
    public User(...) {...}
    public String[] getForwardedAddresses() {...}

    public void setForwardedAddresses(String[] f) {...}
}

public class ConfigurationManager {
    private User loadUser(...) {
        ...
        User user = new User(...);
        String[] f = ...;

        user.setForwardedAddresses(f);
        return user;
    }
}
```

Example of an update (JavaEmailServer)

```
public class User {
    private final String username, domain, password;
-   private String[] forwardAddresses;
+   private EmailAddress[] forwardAddresses;
    public User(...) {...}
-   public String[] getForwardedAddresses() {...}
+   public EmailAddress[] getForwardedAddresses() {...}
-   public void setForwardedAddresses(String[] f) {...}
+   public void setForwardedAddresses(EmailAddress[] f) {...}
}

public class ConfigurationManager {
    private User loadUser(...) {
        ...
        User user = new User(...);
-       String[] f = ...;
+       EmailAddress[] f = ...;
        user.setForwardedAddresses(f);
        return user;
    }
}
```

Example of an update (JavaEmailServer)

```
public class v131_User {
    private final String username, domain, password;
    private String[] forwardAddresses;
}

public class JvolveTransformers {
    ...
    public static void jvolveClass(User unused) {}
    public static void jvolveObject(User to, v131_User from) {
        to.username = from.username;
        to.domain = from.domain;
        to.password = from.password;
        // to.forwardAddresses = null;
        int len = from.forwardAddresses.length;
        to.forwardAddresses = new EmailAddress[len],
        for (int i = 0; i < len; i++) {
            to.forwardAddresses[i] =
                new EmailAddress(from.forwardAddresses[i]);
        }
    }
}
```

Stub generated by UPT
for the old version

Default transformer copies
old fields, initializes new
ones to null

Example of an update (JavaEmailServer)

```
public class v131_User {
    private final String username, domain, password;
    private String[] forwardAddresses;
}

public class JvolveTransformers {
    ...
    public static void jvolveClass(User unused) {}
    public static void jvolveObject(User to, v131_User from) {
        to.username = from.username;
        to.domain = from.domain;
        to.password = from.password;
        // to.forwardAddresses = null;
        int len = from.forwardAddresses.length;
        to.forwardAddresses = new EmailAddress[len];
        for (int i = 0; i < len; i++) {
            to.forwardAddresses[i] =
                new EmailAddress(from.forwardAddresses[i]);
        }
    }
}
```

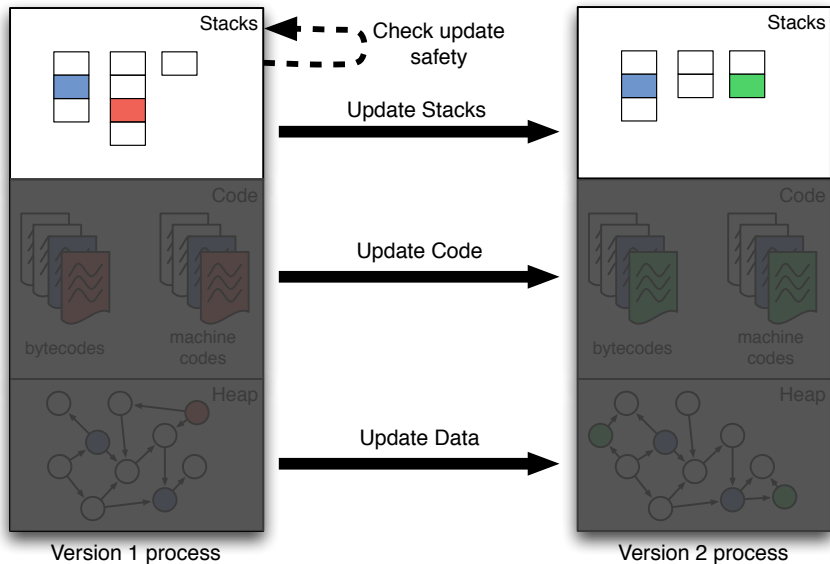
Stub generated by UPT
for the old version

Compiling transformation functions

```
public static void jvolveObject(User to, v131_User from) {  
    to.username = from.username;  
    ...  
}
```


- Very close to Java semantics
- Field username is private and final
- Functions compiled specially by a JastAddJ extension to the Java language
- Ignores access protection and allows assigning to final fields

Check for update safety




Safe point for the update

- Update must be atomic
- Updates happen at “safe points”
- Safe points are VM yield points, and restrict what methods can be on stack
- Extend the thread scheduler to suspend all application threads
- If any stack has a restricted method, delay the update




Safe point for the update

- Update must be atomic
- Updates happen at “safe points”
- Safe points are VM yield points, and restrict what methods can be on stack
- Extend the thread scheduler to suspend all application threads
- If any stack has a restricted method, delay the update



Safe point for the update

- Update must be atomic
- Updates happen at “safe points”
- Safe points are VM yield points, and restrict what methods can be on stack
- Extend the thread scheduler to suspend all application threads
- If any stack has a restricted method, delay the update



Restricted methods

- (1) Methods changed by the update
- (2) Methods identified by the user as unsafe based on semantic information about the application

Install return barriers that trigger DSU upon unsafe method's return

- (3) Methods whose bytecode is unchanged, but compiled representation is changed by the update
 - Offsets of fields and methods hard-coded in machine code
 - Inlined callees may have changed

Utilize on-stack replacement to recompile base-compiled methods

Restricted methods


- (1) Methods changed by the update
- (2) Methods identified by the user as unsafe based on semantic information about the application

Install return barriers that trigger DSU upon unsafe method's return

- (3) Methods whose bytecode is unchanged, but compiled representation is changed by the update
 - Offsets of fields and methods hard-coded in machine code
 - Inlined callees may have changed


Utilize on-stack replacement to recompile base-compiled methods

Reaching a safe point



Install a return barrier for d(). Wait till it returns. On-stack replace new machine code for c().

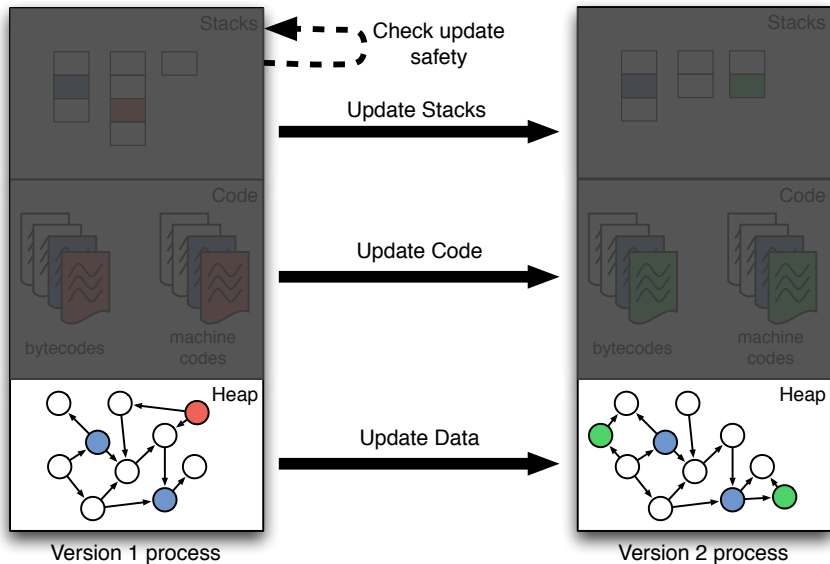
Update code




Update code

- Modify class loader to recognize new versions of classes
- Install new versions of classes and methods
- Rely on Just-in-time Compiler to compile new versions of methods on demand
- Extend On-stack replacement to update active methods

Update data




Transforming objects in the GC




Happens in two steps

- Garbage collector creates an additional empty copy for updated objects
- Walk through and transform all these objects


Jvolve GC



Jvolve GC




Jvolve GC




```
jvolveObject(Node to,
               old_Node from) {
    to.data = from.data;
    to.next = from.next;
    if (to.next != null)
        to.next.prev = to;
}
```

Jvolve GC




```
jvolveObject(Node to,  
               old_Node from) {  
    to.data = from.data;  
    to.next = from.next;  
    if (to.next != null)  
        to.next.prev = to;  
}
```

Jvolve GC




```
jvolveObject(Node to,
               old_Node from) {
    to.data = from.data;
    to.next = from.next;
    if (to.next != null)
        to.next.prev = to;
}
```

Jvolve GC




```
jvolveObject(Node to,
               old_Node from) {
    to.data = from.data;
    to.next = from.next;
    if (to.next != null)
        to.next.prev = to;
}
```


Jvolve GC




Jvolve GC




Jvolve GC



Jvive GC



Jvolve GC



Outline

- Introduction
- Jvolve
 - VM Implementation
 - Evaluation
- Future Work
- Conclusion

Application Experience

- Jetty webserver
 - 11 versions, 5.1.0 through 5.1.10, 1.5 years
 - 45 KLOC
- JavaEmailServer
 - 10 versions, 1.2.1 through 1.4, 2 years
 - 4 KLOC
- CrossFTP server
 - 4 versions, 1.05 through 1.08, more than a year
 - 18 KLOC

Support 20 of 22 updates

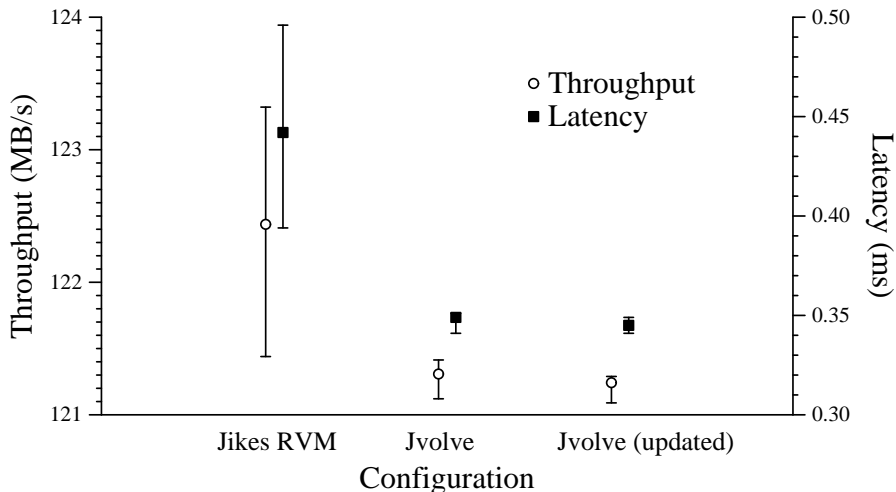
- 13 updates change class signature by adding new fields
- Several updates require On-stack replacement support
- Two versions update an infinite loop, postponing the update indefinitely

Unsupported updates

- JavaEmailServer 1.2.4 to 1.3
 - Update reworks the configuration framework of the server
 - Many classes are modified to refer to the configuration system
 - Including infinite loops in SMTP and POP threads
- Jetty 5.1.2 to 5.1.3
 - The application would never reach a safe point
 - Modified method `ThreadedServer.acceptSocket()` that waits for connections is nearly always on stack
 - Return barrier not sufficient since the main method in other threads `PoolThread.run()` is itself modified

JVOLVE performance


No overhead during steady-state execution



Update pause time

- No apriori overhead during normal execution (before and after the update)
- Only effect on execution time is the update pause time
 - Comparable to GC pause time

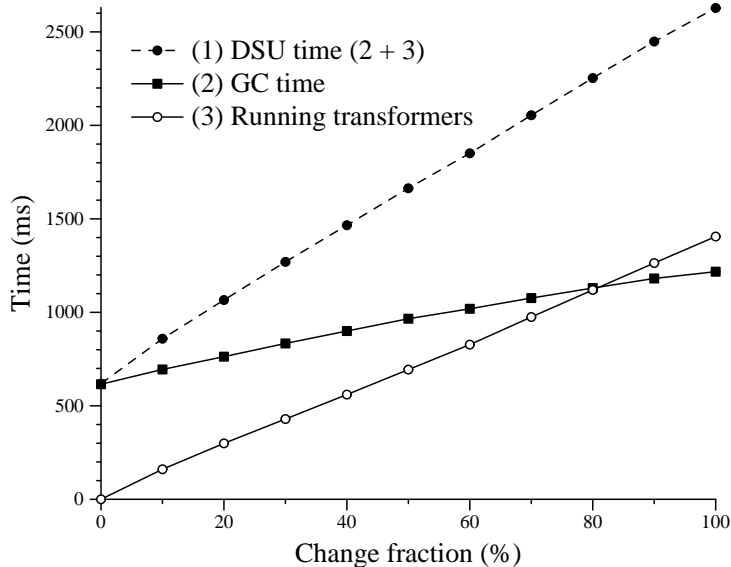
Update pause time



Update pause time

$$\begin{aligned} \text{DSU Pause Time} &\cong \text{Regular GC Time} + \\ &\quad \text{Time to allocate upd. objects} + \\ &\quad \text{Time to transform objects} \\ &\propto \text{Upd. objects fraction} \\ &\quad \text{Heap size} \end{aligned}$$

DSU pause times (microbenchmark)



Outline

- Introduction
- Jvolve
 - VM Implementation
 - Evaluation
- Future Work
- Conclusion

What's in the path towards
mainstream adoption?

- Currently, we guarantee type safety. Can we do more?
- Safe behavior in a general context is undecidable
[Gupta et al., 1996]
- Divide program into transactions, ensure they are version consistent [Neamtiu et al., 2008]
- Exhaustively testing updates

- Supporting every possible update is not a goal
- Support most of the updates that occur in practice
- Large and complex applications and updates
- Understand update semantics from refactorings

- Update time roughly proportional to GC time
- Semi-space GC is not practical
- We have two requirements
 - Copying objects
 - Full-heap collection
- Other collectors? Concurrent collectors?

Conclusion



- Jvolve, a Java VM with support for Dynamic Software Updating
- Most-featured, best-performing DSU system for Java
- Naturally extends existing VM services
- Supports about two years worth of updates

Dynamic software updating in managed languages can be achieved in a safe, flexible and efficient manner.

Source code and other information:

<http://www.cs.utexas.edu/~suriya/jvolve>

References I

-  Gupta, D., Jalote, P., and Barua, G. (1996).
A formal framework for on-line software version change.
IEEE Trans. Softw. Eng., 22(2):120–131.
-  Neamtiu, I., Hicks, M., Foster, J. S., and Pratikakis, P. (2008).
Contextual effects for version-consistent dynamic software
updating and safe concurrent programming.
In *Proc. POPL*.