

# Dynamic Software Updates: A VM-centric Approach

Suriya Subramanian<sup>1</sup>   Michael Hicks<sup>2</sup>  
Kathryn S. McKinley<sup>1</sup>

<sup>1</sup>Department of Computer Sciences  
The University of Texas at Austin

<sup>2</sup>Department of Computer Science  
University of Maryland

April 17, 2009  
Parasol Seminar Series



*The only thing that is constant is change.*

*Heraclitus of Ephesus*

# Motivation

- Software applications change all the time
- Deployed systems must be updated with bug fixes, new features
- The straightforward approach is to stop and restart applications
- Stopping not desirable
  - Safety concerns
  - Revenue loss
  - Inconvenience

# Applications

- Personal operating system
- High availability enterprise applications
- Even a cache with lots of state

---

<sup>1</sup><http://hurvitz.org/blog/2008/06/linkedin-architecture>

# Applications

- Personal operating system
- High availability enterprise applications
- Even a cache with lots of state
  - LinkedIn.com architecture<sup>1</sup>
    - “The Cloud”: In memory representation of the LinkedIn network graph
    - Network size - 22M nodes, 120M edges
    - Rebuilding an instance takes 8 hours

---

<sup>1</sup><http://hurvitz.org/blog/2008/06/linkedin-architecture>

# Solutions to updating software

- Move state out of the process
  - State stored externally, for instance databases
  - Redundant systems: start a new process and stop this one
  - Not always possible
- Dynamic Software Updating (DSU)
  - Update process state without restarting application
  - Non-redundant systems benefit as well
  - Decouples fault-tolerance from software updating

# Solutions to updating software

- Move state out of the process
  - State stored externally, for instance databases
  - Redundant systems: start a new process and stop this one
  - Not always possible
- Dynamic Software Updating (DSU)
  - Update process state without restarting application
  - Non-redundant systems benefit as well
  - Decouples fault-tolerance from software updating

# DSU requirements

A Dynamic Software Updating solution should *ideally* be

**Safe** Updating is as correct as starting from scratch

**Flexible** Be able to support changes encountered in practice

**Efficient** No performance impact on the original application



# State of the art

## Significant progress for C

- Server feature upgrades
  - Ginseng [Neamtiu et al., 2006]
  - POLUS [Chen et al., 2007]
- Security patches: OPUS [Altekar et al., 2005]
- Operating system upgrades
  - K42 [Soules et al., 2003]
  - DynAMOS [Makris and Ryu, 2007]
  - LUCOS [Chen et al., 2006]
  - Ksplice [Arnold and Kaashoek, 2009]

# Opportunities for managed languages

## Solutions for C typically

- Require special compilation
- Statically/dynamically insert indirection for function calls
- Restrict structure updates, require extra allocation
- Impose space/time overheads on normal execution
- Make type-safety for updates difficult
- Not multi-threaded

# Existing solutions for managed languages

- VM-based solutions
  - JDrums [Ritzau and Andersson, 2000], DVM [Malabarba et al., 2000]
  - Not well evaluated
  - Provide an interface similar to JVOLVE
  - Perform lazy updates
  - Overheads during normal execution
- Standard VM with DSU support
  - DJVCS [Barr and Eisenbach, 2003], DUSC [Orso et al., 2002], [Milazzo et al., 2005]
  - Special classloaders, compilers
  - Very restrictive
  - Space/time overheads

# Our solution

- JVOLVE - a Java Virtual Machine with DSU support
- Key insight: Naturally extend existing VM services
  - Classloading
  - Bytecode verification<sup>2</sup>
  - Thread synchronization
  - JIT Compilation
  - On-stack replacement
  - Garbage collection
- No DSU-related overhead during normal execution
- Support updates to real world applications

---

<sup>2</sup>Jikes RVM does not have a bytecode verifier

# Claim

*Dynamic software updating in managed languages can be achieved in a safe, flexible and efficient manner by naturally extending existing VM services.*

# Claim

*Dynamic software updating in managed languages can be achieved in a safe, flexible and efficient manner by naturally extending existing VM services.*

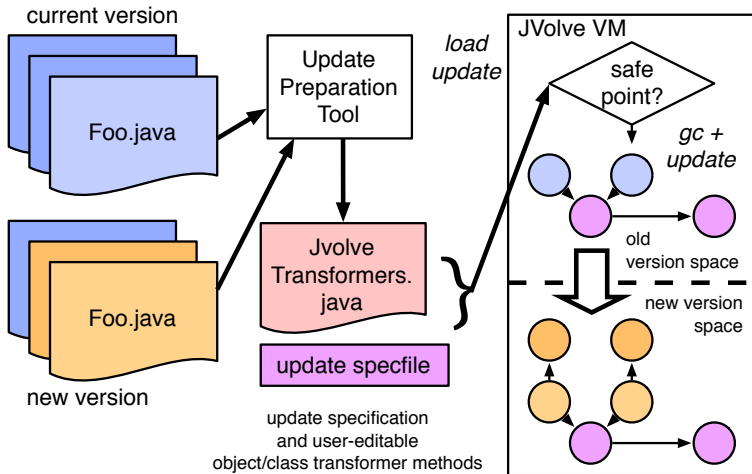
## Corollary

*DSU support should be a standard feature of future VMs.*

# Outline

- Introduction
  - Motivation
  - Solutions
- JVOLVE
  - Developer's view
  - Implementation
  - Experience
- Conclusion

# Developer's view of JVOLVE





# Division of Labor

- Developer
  - Write the old and new versions
  - Write class/object transformation functions for classes that changed (optional)
  - Testing (both the application and the update)
- JVOLVE system
  - Update Preparation Tool (UPT) compares versions and presents the update to the JVOLVE VM.
  - JVOLVE VM handles the update

# Supported updates

- Changes within the body of a method
- Class signature updates
  - Add, remove, change the type signature of fields and methods
- Changes can occur at any level of the class hierarchy

# Example of an update (JavaEmailServer)

```
public class User {
    private final String username, domain, password;
    private String[] forwardAddresses;
    public User(...) {...}
    public String[] getForwardedAddresses() {...}
    public void setForwardedAddresses(String[] f) {...}
}

public class ConfigurationManager {
    private User loadUser(...) {
        ...
        User user = new User(...);
        String[] f = ...;
        user.setForwardedAddresses(f);
        return user;
    }
}
```

# Example of an update (JavaEmailServer)

```
public class User {
    private final String username, domain, password;
-   private String[] forwardAddresses;
-   public String[] getForwardedAddresses() {...}
    public User(...) ...
-   public void setForwardedAddresses(String[] f) {...}
+   private EmailAddress[] forwardAddresses;
+   public EmailAddress[] getForwardedAddresses() {...}
+   public void setForwardedAddresses(EmailAddress[] f) {...}
}

public class ConfigurationManager {
    private User loadUser(...) {
        ...
        User user = new User(...);
-       String[] f = ...;
+       EmailAddress[] f = ...;
        user.setForwardedAddresses(f);
        return user;
    }
}
```

# Object Transformers

- “Transform” objects to correspond to the new version
- A function generated by the Update Preparation Tool (UPT)
- Accepts old object and new object as parameters
- Default transformer copies old fields and initializes new ones to `null`
- User can optionally modify this function

# Object Transformers

```
public class v131_User {  
    private final String username, domain, password;  
    private String[] forwardAddresses;  
}  
  
public class JvolveTransformers {  
    ...  
    public static void jvolveClass(User unused) {}  
    public static void jvolveObject(User to, v131_User from) {  
        to.username = from.username;  
        to.domain = from.domain;  
        to.password = from.password;  
        // to.forwardAddresses = null;  
        int len = from.forwardAddresses.length;  
        to.forwardAddresses = new EmailAddress[len];  
        for (int i = 0; i < len; i++) {  
            String[] parts = from.forwardAddresses[i].split("@", 2);  
            to.forwardAddresses[i] = new EmailAddress(parts[0], parts[1]);  
        }  
    }  
}
```

Stub generated by UPT for  
the old version

Default transformer copies  
old fields, initializes new  
ones to null

# Object Transformers

```
public class v131_User {  
    private final String username, domain, password;  
    private String[] forwardAddresses;  
}  
  
public class JvolveTransformers {  
    ...  
    public static void jvolveClass(User unused) {}  
    public static void jvolveObject(User to, v131_User from) {  
        to.username = from.username;  
        to.domain = from.domain;  
        to.password = from.password;  
        // to.forwardAddresses = null;  
        int len = from.forwardAddresses.length;  
        to.forwardAddresses = new EmailAddress[len];  
        for (int i = 0; i < len; i++) {  
            String[] parts = from.forwardAddresses[i].split("@", 2);  
            to.forwardAddresses[i] = new EmailAddress(parts[0], parts[1]);  
        }  
    }  
}
```

Stub generated by UPT for  
the old version

# Outline

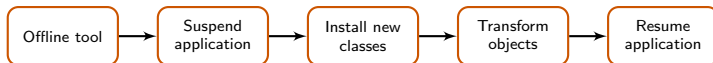
- Introduction
  - Motivation
  - Solutions
- **JVOLVE**
  - Developer's view
  - **Implementation**
  - Experience
- Conclusion



# Update model


- Update happens in one fell swoop
- Simple to reason about
- Code
  - Old code before the update
  - New code after the update
- Data
  - Representation consistency (all values of a type correspond to the latest version)
  - Support a transformation function to convert objects to conform to their new definition

# Update process



- Offline Update Preparation Tool (UPT)
- JVM VM
  - Reach a safe point in the VM (thread synchronization)
  - Install new classes (classloader)
  - Transform objects to new definition (garbage collector)
  - Resume execution

# Update Preparation Tool



- Uses jclasslib<sup>3</sup>, a bytecode library
- Compares bytecodes of the two versions
- Categorizes changes into
  - Updated classes** Classes that add, remove, change signature of fields or methods
  - Updated methods** Changes within a method body. Only the method has to be loaded/updated
  - Indirect updates** No change to method, but refers to changed classes
- Generates old version stubs and default object transformers

---

<sup>3</sup><http://jclasslib.sourceforge.net>

# Compiling transformation functions

- All transformers specified in a separate source file

- Class transformers

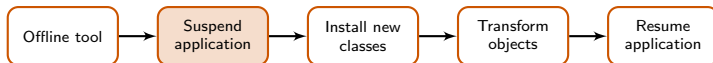
```
jvolveClass(ClassName unused)
```

- Object transformers

```
jvolveObject(old_ClassName from, ClassName to)
```

- Compiled specially by a JastAddJ extension to the Java language
- Ignores access protection and allows assigning to `final` fields

# Safe point for the update



- Update must be atomic
- Updates happen at “safe points” (VM yield points with restriction on what methods can be on stack)
- Extend the thread scheduler to suspend all application threads
- Examine all stacks, ensure no restricted methods on stack and perform the update

# Restricted methods

- (1) Methods changed by the update
- (2) Methods whose bytecode is unchanged, but compiled representation is changed by the update
  - Offsets of fields and methods hard-coded in machine code
  - Inlined callees may have changed
- (3) Methods identified by the user as unsafe based on semantic information about the application

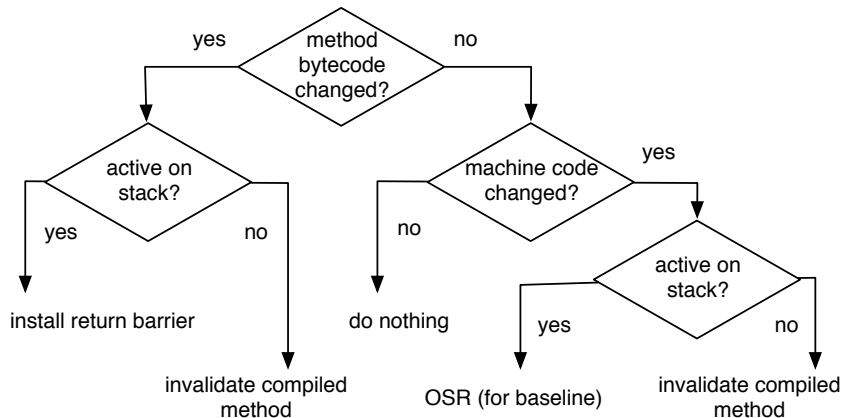
# Restricted methods

- (1) Methods changed by the update
- (2) Methods whose bytecode is unchanged, but compiled representation is changed by the update
  - Offsets of fields and methods hard-coded in machine code
  - Inlined callees may have changed
- (3) Methods identified by the user as unsafe based on semantic information about the application

## Handling restricted methods

- *On-stack replace* baseline-compiled category (2) methods
- Do not allow (1) and (3) to be active on stack, install a return barrier for such methods

# Handling restricted methods





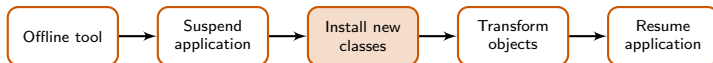
# On stack replacement in JVOLVE

- Used in Jikes RVM to optimize long running methods
- JVOLVE utilizes OSR for DSU
- Currently only support baseline-compiled methods
- Can OSR any method on stack
- Extract the state of the stack
- Construct a new method with a specialized prologue (at the bytecode level) that reconstructs the stack
- Last instruction of prologue jumps to bytecode where execution should resume from
- Overwrite the return address to point to the special method

# On stack replacement in JVOLVE

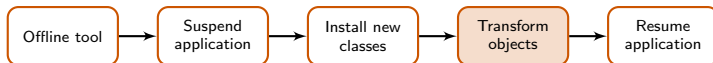
- Used in Jikes RVM to optimize long running methods
- JVOLVE utilizes OSR for DSU
- Currently only support baseline-compiled methods
- Can OSR any method on stack
- Extract the state of the stack
- Construct a new method with a specialized prologue (at the bytecode level) that reconstructs the stack
- Last instruction of prologue jumps to bytecode where execution should resume from
- Overwrite the return address to point to the special method

# Installing new classes





- The VM maintains Class, Method and Field data structures
- For Method updates: Only load the new method's bytecodes
- For Class updates: Rename the old class and load the entire class file (equivalent to have loaded two different class)

# Transforming objects




- Built on top of a semi-space copying collector
- As part of collector's visit allocate additional space for updated objects
- After GC, run class and object transformers

# Semi-space copying collector




The heap is divided into two spaces. Only one space is used by the application. The garbage collector copies objects from *FromSpace* to *ToSpace*.

# Semi-space copying collector




GC copies *A* to *ToSpace*, leaves a forwarding pointer pointing to the new copy *A'*.

# Semi-space copying collector



GC scans A'. The objects pointed to by A' (B and C) are copied to *ToSpace*. A's fields point to the copied objects.


# Semi-space copying collector



Next, the GC scans B', and copies objects D and E.





# Semi-space copying collector





Similarly for C'

# Semi-space copying collector




Similarly for C', D'

# Semi-space copying collector




Similarly for C', D', and E.

# Semi-space copying collector



When scanning F', the first field points to A in *FromSpace*, which is a forwarding pointer. After the scan, this field points to A'.

# Semi-space copying collector




All objects in *ToSpace* are scanned. All reachable/live objects are now in *ToSpace*.

# JVOLVE Garbage collector


- Identical to Semispace for “regular” objects
- For objects to be transformed
  - Copy the object to ToSpace (like Semispace)
  - Also, allocate an empty object in ToSpace for the new version
- Forwarding pointers point to the “new version” object
- No field can point to an “old version” object



# JVOLVE garbage collector



The same heap as before. Objects to be transformed are highlighted.

# JVOLVE garbage collector




 To be transformed  
  $v_1$  object

Copy A.

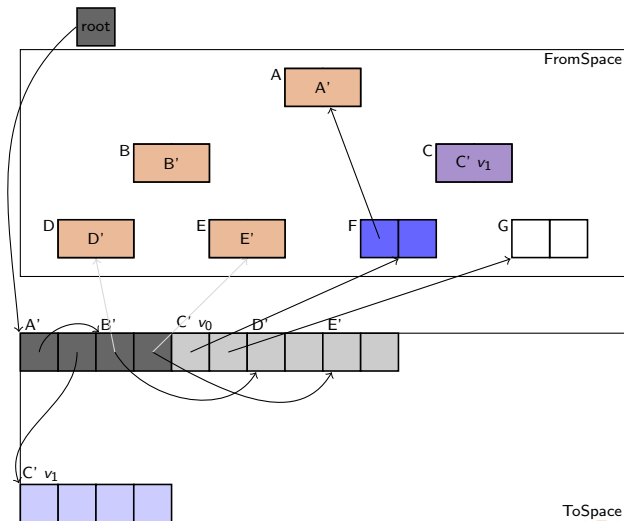


# JVOLVE garbage collector




Scan A'. Copy B and C. In addition an empty object C'v<sub>1</sub> is allocated. A' points to this copy and not the old one.


# JVOLVE garbage collector



# JOLVE garbage collector




# JOLVE garbage collector



Scan D'.

# JOLVE garbage collector

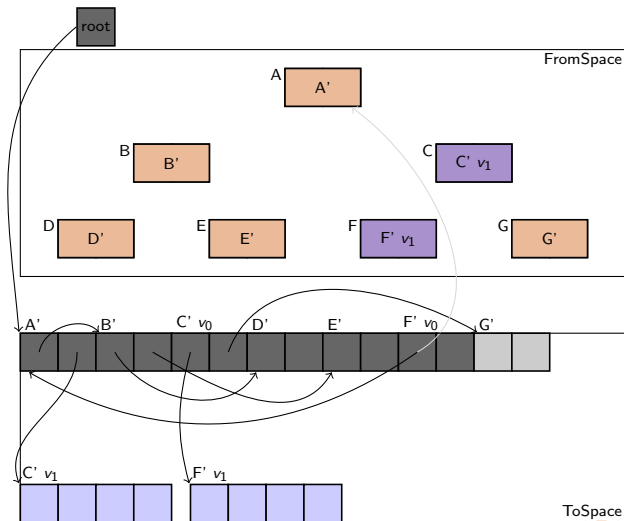


To be transformed

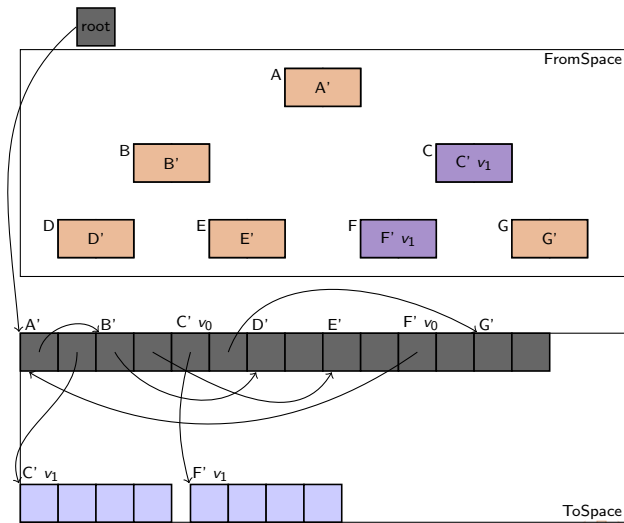
$v_1$  object

Scan E'.

# JOLVE garbage collector




# JVOLVE garbage collector



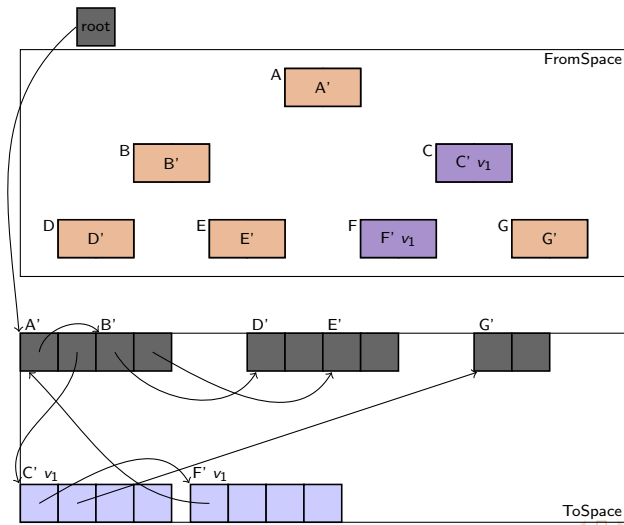
GC is now complete.  
No field can point to  $C' v_0$  or  $F' v_0$ . Pointers to C and F point to  $v_1$  (empty) objects.  
`memcpy(v_1, v_0);`  
will give us a valid heap.

# JVOLVE garbage collector






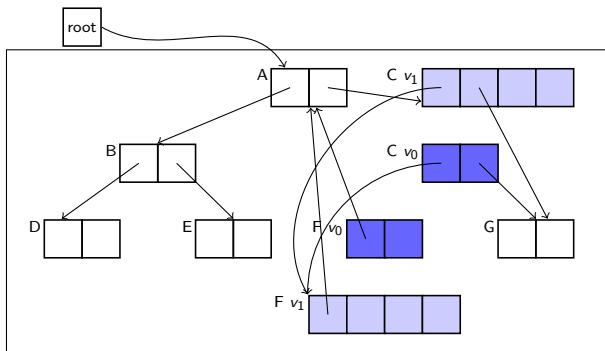
# JVOLVE garbage collector




# JVOLVE Garbage collector




# JOLVE Garbage collector



# JVOLVE Garbage collector



# Revisiting transformation functions



We have an ordering problem

`(C v0).field0.field0` might be uninitialized

# Revisiting transformation functions

## Solutions to the ordering problem

- Programmer can invoke a VM function that will transform objects on demand. Moves burden of safety to the programmer
- Insert read barrier code to perform this check when compiling the transformation function
- Perform some static analysis to determine an order to queue objects

# Revisiting transformation functions

## Solutions to the ordering problem

- Programmer can invoke a VM function that will transform objects on demand. Moves burden of safety to the programmer
- Insert read barrier code to perform this check when compiling the transformation function
- Perform some static analysis to determine an order to queue objects

# Updating from a singly to doubly linked list

```
public static void jvolveObject(LinkedList.Node to,  
                                r0_LinkedList.Node from) {  
    to.next = from.next;  
    to.data = from.data;  
    if (to.next != null) to.next.prev = to;  
}
```

```
public static void jvolveObject(LinkedList to, r0_LinkedList from) {  
    Jvolve.transformReferences(from);  
    to.head = from.head;  
    LinkedList.Node n0 = null;  
    LinkedList.Node n1 = to.head;  
    while (n1 != null) {  
        n0 = n1;  
        n1 = n1.next;  
    }  
    to.tail = n0;  
}
```



# Outline

- Introduction
  - Motivation
  - Solutions
- JVOLVE
  - Developer's view
  - Implementation
  - Experience
- Conclusion

# Applications

- Jetty webserver
  - 11 versions, 5.1.0 through 5.1.10, 1.5 years
  - 45 KLOC
- JavaEmailServer
  - 10 versions, 1.2.1 through 1.4, 2 years
  - 4 KLOC
- CrossFTP server
  - 4 versions, 1.05 through 1.08, more than a year
  - 18 KLOC


# Overhead of DSU

- No discernible overhead for normal execution (before and after the update)
- Only effect on execution time is the update pause time
  - Comparable to GC pause time

# Jetty webserver performance

- Used `httperf` to issue requests
- Create 800 new connections/second (saturation rate)
- 5 serial requests to 40KB file per connection
- Compared versions 5.1.5 and 5.1.6
- Experiments on Intel Core 2 Quad, Linux 2.6.22, JikesRVM  
SVN r15532


# Jetty webserver: Throughput measurements



# DSU pause times

- JVOLVE performs a GC to transform objects
- Pause time determined by
  - Heap size
  - # of objects transformed
- Simple microbenchmark varying the # of objects transformed

# DSU pause times (microbenchmark)



# Jetty webserver: Summary of changes

Ver.	# classes added	classes	# changed methods			fields	
			add	del	chg	add	del
5.1.1	0	14	4	1	38/0	0	0
5.1.2	1	5	0	0	12/1	0	0
5.1.3	3	15	19	2	59/0	10	1
5.1.4	0	6	0	4	9/6	0	2
5.1.5	0	54	21	4	112/8	5	0
5.1.6	0	4	0	0	20/0	5	6
5.1.7	0	7	8	0	11/2	9	3
5.1.8	0	1	0	0	1/0	0	0
5.1.9	0	1	0	0	1/0	0	0
5.1.10	0	4	0	0	4/0	0	0



# Unsupported updates

- Jetty 5.1.2 to 5.1.3
  - The application would never reach a safe point
  - Modified method `ThreadedServer.acceptSocket()` that waits for connections is nearly always on stack
  - Return barrier not sufficient since the main method in other threads `PoolThread.run()` is itself modified
- JavaEmailServer 1.2.4 to 1.3
  - Update reworks the configuration framework of the server
  - Many classes are modified to refer to the configuration system
  - Including infinite loops in SMTP and POP threads

# Outline

- Introduction
  - Motivation
  - Solutions
- JVOLVE
  - Developer's view
  - Implementation
  - Experience
- Conclusion




# Conclusion

- JVOLVE, a Java VM with support for Dynamic Software Updating
- Most-featured, best-performing DSU system for Java
- Extends existing VM services
- Supports about two years worth of updates




*Dynamic software updating in managed languages can be achieved in a safe, flexible and efficient manner by naturally extending existing VM services.*

# Thank you

# References I

-  Altekar, G., Bagra, I., Burstein, P., and Schultz, A. (2005).  
OPUS: Online patches and updates for security.  
*In Proc. USENIX Security.*
-  Arnold, J. and Kaashoek, F. (2009).  
Ksplice: Automatic rebootless kernel updates.  
*In Proc. EuroSys.*
-  Barr, M. and Eisenbach, S. (2003).  
Safe upgrading without restarting.  
*In Proc. ICSM.*

# References II

-  Chen, H., Chen, R., Zhang, F., Zang, B., and Yew, P.-C. (2006).  
Live updating operating systems using virtualization.  
*In Proc. VEE.*
-  Chen, H., Yu, J., Chen, R., Zang, B., and Yew, P.-C. (2007).  
POLUS: A POverful Live Updating System.  
*In Proc. ICSE.*
-  Makris, K. and Ryu, K. D. (2007).  
Dynamic and adaptive updates of non-quiescent subsystems in  
commodity operating system kernels.  
*In Proc. EuroSys.*

## References III



Malabarba, S., Pandey, R., Gragg, J., Barr, E., and Barnes, J. F. (2000).

Runtime support for type-safe dynamic Java classes.  
*In Proc. ECOOP.*






Milazzo, M., Pappalardo, G., Tramontana, E., and Ursino, G. (2005).

Handling run-time updates in distributed applications.  
*In Proc. SAC.*



Neamtiu, I., Hicks, M., Stoye, G., and Oriol, M. (2006).  
Practical dynamic software updating for C.  
*In Proc. PLDI.*

# References IV

-  Orso, A., Rao, A., and Harrold, M. J. (2002).  
A technique for dynamic updating of Java software.  
*In Proc. ICSM.*
-  Ritzau, T. and Andersson, J. (2000).  
Dynamic deployment of Java applications.  
*In Proc. Java for Embedded Systems Workshop.*
-  Soules, C., Appavoo, J., Hui, K., Silva, D. D., Ganger, G., Krieger, O., Stumm, M., Wisniewski, R., Auslander, M., Ostrowski, M., Rosenberg, B., and Xenidis, J. (2003).  
System support for online reconfiguration.  
*In Proc. USENIX Annual Technical Conference.*